

Decision Trees and Random Forests

Data science Certificate

Aymeric DIEULEVEUT

May 2022

1 Decision Trees

2 Random Forests

Outline

1 Decision Trees

2 Random Forests

A Reminder on Supervised Learning

- Observations: $(X_i)_{i=1,\dots,n}$. Each X_i has d components.
- Outputs: $(Y_i)_{i=1,\dots,n}$

3 Examples:

- Iris Dataset:
- MNIST:
- House price prediction

Example: Iris Dataset



Figure: Iris: setosa, versicolor, virginica

One of the most widely used toy dataset in classification:

- 3 classes : [*'setosa'*, *'versicolor'*, *'virginica'*]
- 50 points per class.
- 4 features: [*'sepal length'*, *'sepal width'*, *'petal length'*, *'petal width'*]

Example: MNIST Dataset




Figure:

Digit recognition from an image:

- 10 classes : $Y_i \in$
- 60k images, (50k train, 10k test), balanced
- 784 features.

House price Prediction

- Output : house value $Y \in$
- 1500 examples
- 81 features: Construction year, Neighborhood, Surface, Floor, Balcony: both quantitative and qualitative, some ordinal variables.

Summary of those three datasets

Iris

MNIST

House price prediction

Number of examples n

Number of the features d

Type of the features

Space \mathcal{Y}

Task

Decision Tree

Goals of Decision Trees

- ➊ Supervised Learning: solve the classification or regression task
- ➋ Provide some understanding: what on this example allows me to classify it? What are the important features ?

Decision Tree

Goals of Decision Trees


- ➊ Supervised Learning: solve the classification or regression task
- ➋ Provide some understanding: what on this example allows me to classify it? What are the important features ?

Two questions:

- How to train a decision tree? → **training phase (learning algorithm)**
- How to predict with a given decision tree? → **inference phase**

Inference Phase

- Root = observation = input
- Nodes = tests
- Branches = possible outcomes
- Leaves = final decision = output




Training: how to build a tree.

Key question !


- What is a good tree?
- What do I need to choose?
- How to choose?

What do you think?

Building trees : CART




Building trees : CART



How to chose the best cut ?

- e.g., Given a possible cut, measure the reduction of impurity:
 - ▶ in regression: mean-squared-error $\sum_{i \in C} (Y_i - \bar{Y}_C)^2$
 - ▶ in classification: Gini's impurity.
- Find dimension and threshold with optimal impurity reduction.
- Iterate until stopping criterion is met.

Gini's impurity



Gini impurity measures **how often a randomly chosen element from the set** would be **incorrectly labeled if it was randomly labeled according to the distribution of labels in the subset**.

$$I_G(p) = \sum_{i=1}^J p_i(1 - p_i) = 1 - \sum_{i=1}^J p_i^2$$

For two classes, related to the variance if classified as Bernoulli r.v..


Example: Iris Dataset




Figure: Iris: setosa, versicolor, virginica

One of the most widely used toy dataset in classification:

- 3 classes : [*'setosa'*, *'versicolor'*, *'virginica'*]
- 50 points per class.
- 4 features: [*'sepal length'*, *'sepal width'*, *'petal length'*, *'petal width'*]



Exercise: Guess Classification tree for Iris Dataset



In practice : Scikit Learn

Universal Python library for Machine Learning:

- Very easy to use
- Very well documented
- Open Source

<https://scikit-learn.org/stable/modules/generated/sklearn.tree.DecisionTreeClassifier.html>

Example: Iris Dataset

What do we need to specify ?

```
from sklearn.datasets import load_iris
from sklearn.model_selection import cross_val_score
from sklearn.tree import DecisionTreeClassifier, plot_tree
from sklearn.model_selection import train_test_split

# Choose the Learning algorithm
clf = DecisionTreeClassifier(max_depth = 2,
                             random_state=0)


# Load the dataset
iris = load_iris()
X_train, X_test, y_train, y_test = train_test_split(
    iris.data, iris.target, test_size=0.33, random_state=43)

# Check accuracy
cross_val_score(clf, iris.data, iris.target, cv=10)

# Fit the model
clf.fit(X_train,y_train)

# Plot tree
plot_tree(clf, feature_names = fn,
          class_names=cn,
          filled = True)
plt.show()
```

Output:



Example:

```
# Check
X_test[0,:], iris.target_names[y_test[0]]
```



Returns:

(array([4.8, 3.1, 1.6, 0.2]), 'setosa')

Code: <https://colab.research.google.com/drive/1dXYIjacNAeASg154vIWYwJMQlfgigSVA?usp=sharing#scrollTo=VeJHkJd7RsDw>

Example: Iris Dataset

Output:



Example:

```
# Check
X_test[0,:], iris.target_names[y_test[0]]
```

Returns:

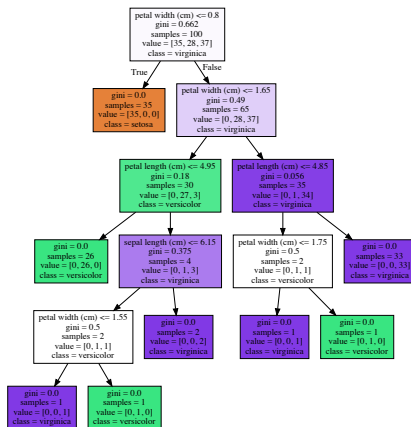
(array([4.8, 3.1, 1.6, 0.2]), 'setosa')

Code: <https://colab.research.google.com/drive/1dXYIjacNAeASgl54vIWYwJMQlfgigSVA?usp=sharing#scrollTo=VeJHkJd7RsDw>

What if I use a deeper tree ? Or if I do not specify the depth in the model definition ?

Deeper trees

```
# Choose the Learning algorithm  
clf = DecisionTreeClassifier(random_state=0)
```



- When does the algorithm stop then ?
- What do you think about it ?

Stopping / Pruning

Trees that achieve perfect classification may suffer from **over-fitting**. (see example in the lab)

To avoid this problem, we can reduce the complexity of the trees:

① Stopping rules

- ▶ Set a minimum number of samples inside each leaf.
- ▶ Set a maximum depth.

② Pruning

- ▶ Reduced error pruning.
- ▶ Cost complexity pruning.

Pruning : example of Cost complexity pruning.

We create a sequence of Trees by changing one subtree at each step into a leaf, until we get only the root: the subtree is chosen to minimize the error made.¹ Then the test error is evaluated along the sequence, to choose the optimal pruning.



¹ *more details here*

Cart: pros and cons

Pros:

- Simple to understand, interpret, visualize
- Implicitly perform features/variable selection
- can handle both categorical and numerical data
- little effort for data preparation
- Non linear relationships do not affect performance
- Can work with large number of observations.

Cons:

- 1 Can create over-complex trees: overfitting
- 2 Unstability: small variations of data can generate completely different trees
- 3 Cannot guarantee global optimal tree
- 4 Biased if some classes dominate

Solution : Using Random Forests !

Outline

1 Decision Trees

2 Random Forests

Random Forests

To put it in simple words: random forests builds multiple decision trees and merges them together to get a more accurate and stable predictions.

→ key idea: create variable trees + aggregate them

How to create variability ?

Random Forests

To put it in simple words: random forests builds multiple decision trees and merges them together to get a more accurate and stable predictions.

→ key idea: create variable trees + aggregate them

How to create variability ?

- Instead of the most important features, search the best feature among a random subset (ntry).
- Work on subsets of data (bootstrap=True).
- More <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.RandomForestRegressor.html>

```
# Creating and fitting a Random Forest
from sklearn.ensemble import RandomForestClassifier

plt.figure(figsize=(20,10))
clf = RandomForestClassifier(n_estimators = 100,
                             max_depth=3, bootstrap = True,
                             random_state =43)

clf.fit(X_train,y_train)
```

How to aggregate ? How to interpret a Random Forest ?

Aggregation

We using a voting or averaging process !

Feature importance

- how much tree nodes using a particular feature reduce impurity along the trees.
- suggests feature selection rule: drop out features with low importance to avoid overfitting.
- Attribute **feature_importance_** in RandomForestRegressor of Sklearn.




Figure: Feature importance output for a Random Forest Classifier in Python

Comparison with decision trees

Cons:

- Less interpretable
- Slower to run

Pros:

- Better in higher dimension
- More stable outputs
- Feature selection

Improving predictions - more arguments.

- ❶ **n_estimators** : number of trees in the forest
 - ▶ improves prediction / stability
 - ▶ slows down the algorithm
- ❷ **max_features** : maximum number of features considered to split a node
- ❸ **min_sample_leaf** : minimum number of samples that should remain in a leaf node.

RF pros and cons

RF Pros:

- 1 works for classification and regression
- 2 default hyperparameters often produce reasonable prediction -> easy to use.
- 3 avoid overfitting if some good trees in the forest and easy feature selection -> high dimensional pb.
- 4 **hard to beat in performance.**
- 5 Bonus : supports multiple outputs!

RF Cons:

- 1 Fast to train but slow to provide new predictions -> ineffective for real-time predictions.
- 2 Good for prediction but bad for description.

What about **Regression** ?

Questions?

Boosting

Another import technique (very powerful !)

Main idea:

- Give more importance to difficult point iteratively
- Incrementally building an ensemble by training each new instance to emphasize the training instances previously mis-modeled.

Example: AdaBoost

See: <https://scikit-learn.org/stable/modules/generated/sklearn.ensemble.AdaBoostClassifier.html>

Bonus: Multiple Outputs

- **Goal:** predict several outputs simultaneously
- **Solution:**
 - ▶ each leaf contains a value for each output
 - ▶ to chose the splits, we use a (weighted) average of the impurity of each output.

Face completion with multi-output estimators



Bonus: complexity

Total complexity for one decision tree:

Worst case:

$$O(n^2 d)$$

Balanced case:

$$O(nd)$$

Tips

- Decision trees tend to overfit: limit the depth or use `min_samples_leaf` (5 is a good initial pick)
- Visualize the tree with a small depth first
- Balance your dataset: trees tend to be biased towards dominating class.

Conclusion

- ❶ Trees are one of the simplest method (the most intuitive / human like)
- ❷ Random Forests give excellent results in many applications.

Lab :

- Example on a synthetic dataset of time series
- Application to inflation prediction in Brazil.

